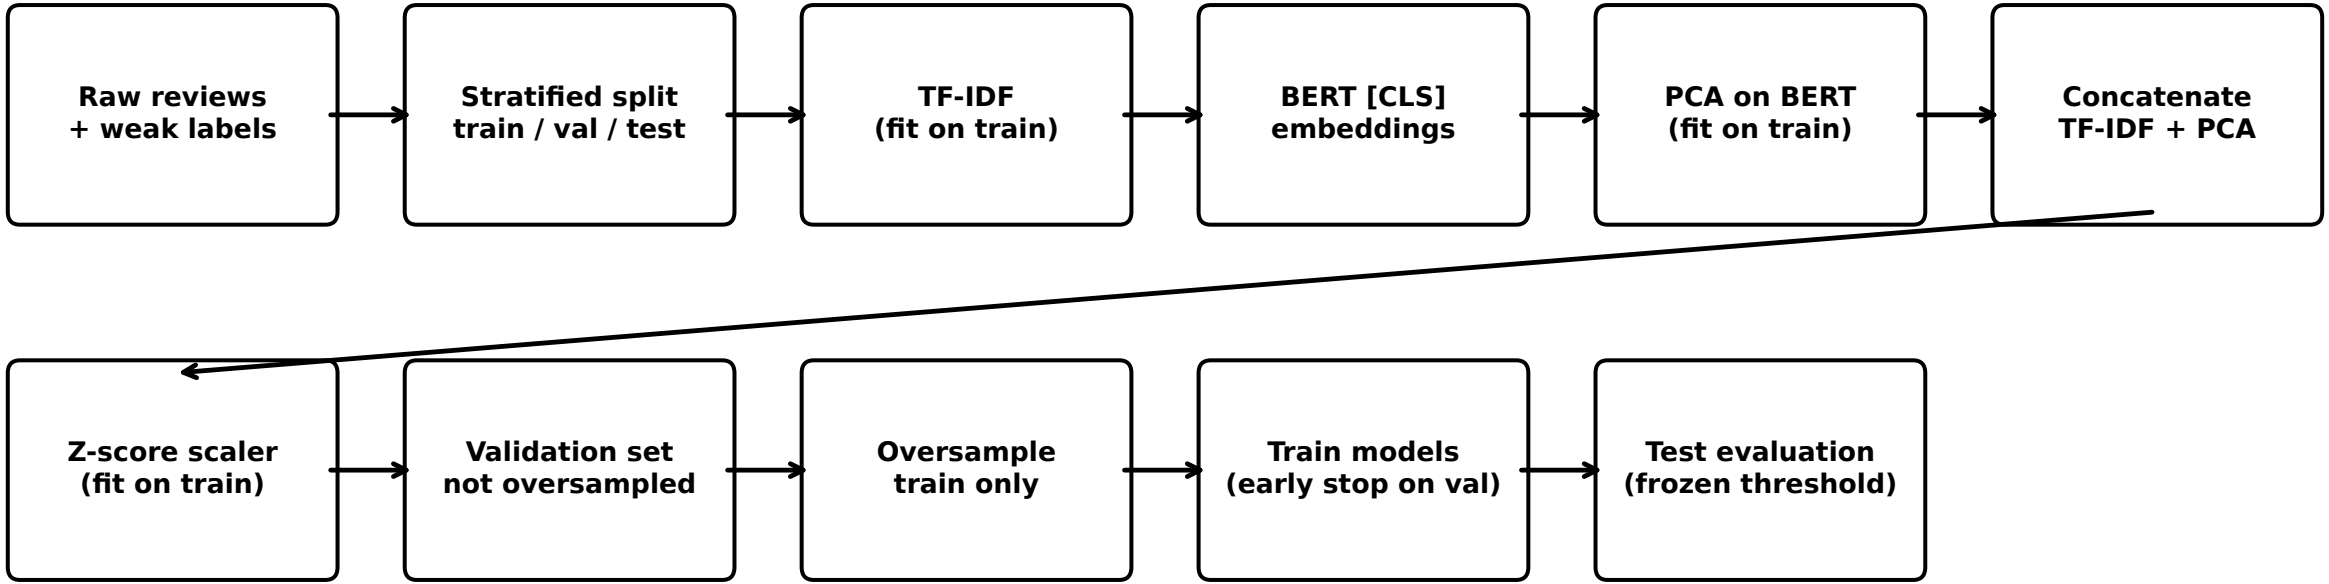
Raw reviews
+ weak labels
Stratified split
train / val / test
TF-IDF
(fit on train)
BERT [CLS]
embeddings
PCA on BERT
(fit on train)
Concatenate
TF-IDF + PCA
Z-score scaler
(fit on train)
Validation set
not oversampled
Oversample
train only
Train models
(early stop on val)
Test evaluation
(frozen threshold)